


A crash course on MPI



Performance of
laptop is only
12-14 years
behind ...
which is about the
time it takes to
complete PhD

Parallel computing: a few things everyone should be familiar with...

- **Amdahl's law:**

S_P – speed-up factor

P - # of processors

ξ - non parallelizable portion of a code

$$S_P = \frac{1}{\xi + (1 - \xi) \frac{1}{P}} < \frac{T_1}{T_P}$$

- **Parallel efficiency**


$$\eta_P = \frac{P-1}{P \log_2 P}$$

- **Communication cost**


$$C = L + \beta l$$

L – latency, l – message size,

β^{-1} - bandwidth



Message Passing Interface - MPI: what do we need it for?



The minimal MPI subset

1. MPI_Init()
2. MPI_Finalize()
3. MPI_Comm_size()
4. MPI_Comm_rank()
5. MPI_Send()
6. MPI_Recv()

```
#include <stdio.h>
#include <mpi.h>

int main (argc, *argv[ ]){
    int rank, size;

    MPI_Init (&argc, &argv);
    /* starts MPI */

    MPI_Comm_rank (MPI_COMM_WORLD, &rank);
    /* get current process id */
    MPI_Comm_size (MPI_COMM_WORLD, &size);
    /* get number of processes */

    printf( "Hello world from process %d of %d\n",
            rank, size );
    MPI_Finalize();
    return 0;
}
```

MPI Communications

- **Point-to-point communications**
 - Involves a sender and a receiver
 - Only the two processors participate in communication
- **Collective communications**
 - All processors within a communicator participate in communication (by calling same routine, may pass different arguments);
 - Barrier, reduction operations, gather, scatter...

Blocking point-to-point communication




```
int MPI_Send(  
    void *buf, /* initial address of send buffer */  
    int count, /* number of elements in send buffer (nonnegative integer) */  
    MPI_Datatype datatype, /* datatype of each send buffer element */  
    int dest, /* rank of destination (integer) */  
    int tag, /* message tag (integer) */  
    MPI_Comm comm /* communicator */  
);
```

```
int MPI_Recv(  
    void *buf, /* initial address of receive buffer */  
    int count, /* number of elements in receive buffer (nonnegative integer) */  
    MPI_Datatype datatype, /* datatype of each receive buffer element */  
    int dest, /* rank of source (integer) */  
    int tag, /* message tag (integer) */  
    MPI_Comm comm, /* communicator */  
    MPI_Status *status /* status object */  
);
```


Deadlock

```
MPI_Comm_rank(MPI_COMM_WORLD,&rank);
If(rank==0)
{
    MPI_Recv(buf1,count,MPI_DOUBLE,1,tag,comm);
    MPI_Send(buf2,count,MPI_DOUBLE,1,tag,comm);
}
else if (rank==1)
{
    MPI_Recv(buf1,count,MPI_DOUBLE,0,tag,comm);
    MPI_Send(buf2,count,MPI_DOUBLE,0,tag,comm);
}
```

```
MPI_Comm_rank(MPI_COMM_WORLD,&rank);
If(rank==0)
{
    MPI_Recv(buf1,count,MPI_DOUBLE,1,tag,comm);
    MPI_Send(buf2,count,MPI_DOUBLE,1,tag,comm);
}
else if (rank==1)
{
    MPI_Send(buf2,count,MPI_DOUBLE,0,tag,comm);
    MPI_Recv(buf1,count,MPI_DOUBLE,0,tag,comm);
}
```



Blocking point-to-point communication



```
int MPI_Sendrecv(  
    void *sendbuf,  
    int sendcount,  
    MPI_Datatype sendtype,  
    int dest,  
    int sendtag,  
  
    void *recvbuf,  
    int recvcount,  
    MPI_Datatype recvtype,  
    int source,  
    int recvtag,  
  
    MPI_Comm comm,  
    MPI_Status *status  
);
```


Example

```
#include <mpi.h>
#include <stdio.h>

int main(int argc, char **argv)
{
    int my_rank, ncpus;
    int left_neighbor, right_neighbor;
    int data_received;
    int send_tag = 101, recv_tag=101;
    MPI_Status status;

    MPI_Init(&argc, &argv);
    MPI_Comm_rank(MPI_COMM_WORLD, &my_rank);
    MPI_Comm_size(MPI_COMM_WORLD, &ncpus);

    left_neighbor = (my_rank-1 + ncpus)%ncpus;
    right_neighbor = (my_rank+1)%ncpus;

    MPI_Sendrecv(&my_rank, 1, MPI_INT, left_neighbor, send_tag,
                 &data_received, 1, MPI_INT, right_neighbor, recv_tag,
                 MPI_COMM_WORLD, &status);

    printf("P%d received from right neighbor: P%d\n",
           my_rank, data_received);

    // clean up
    MPI_Finalize();
    return 0;
}
```

Output:

P3 received from right neighbor: P0

P2 received from right neighbor: P3

P0 received from right neighbor: P1

P1 received from right neighbor: P2

Non-blocking point-to-point communication

```
int MPI_Isend(
    void *buf,      /* initial address of send buffer      */
    int  count,     /* number of elements in send buffer (nonnegative integer) */
    MPI_Datatype datatype, /* datatype of each send buffer element */
    int  dest,      /* rank of destination (integer) */
    int  tag,       /* message tag (integer) */
    MPI_Comm comm,   /* communicator */
    MPI_Request *request /* communication request */
);
```

```
int MPI_Irecv(
    void *buf,      /* initial address of receive buffer */
    int  count,     /* number of elements in receive buffer (nonnegative integer) */
    MPI_Datatype datatype, /* datatype of each receive buffer element */
    int  dest,      /* rank of source (integer) */
    int  tag,       /* message tag (integer) */
    MPI_Comm comm,   /* communicator */
    MPI_Request *request /* communication request */
);
```

What should we use?

MPI_Send + MPI_Recv

MPI_Send + MPI_Irecv


MPI_Isend + MPI_Recv

MPI_Isend + MPI_Irecv

MPI_Sendrecv

MPI_Alltoall

Collective communication



```
int MPI_Reduce (  
    void          *sendbuf,  
    void          *recvbuf,  
    int           count,  
    MPI_Datatype  datatype,  
    MPI_Op        op,  
    int           root,  
    MPI_Comm      comm  
);
```

MPI function


MPI_MAX maximum,
MPI_MIN minimum,
MPI_MAXLOC maximum and location of maximum
MPI_MINLOC minimum and location of minimum
MPI_SUM sum
MPI_PROD product
MPI_LAND logical and
MPI_LOR logical or
MPI_LXOR logical exclusive or
MPI_BAND bitwise and
MPI_BOR bitwise or
MPI_BXOR bitwise exclusive or

Math Meaning


max
min
maximum and location of maximum
minimum and location of minimum
sum
product
logical and
logical or
logical exclusive or
bitwise and
bitwise or
bitwise exclusive or

Implemented in integration, dot products, finding maxima or minima


Collective communication




```
int MPI_Gather (  
    void          *sendbuf,  
    int           sendcnt,  
    MPI_Datatype  sendtype,  
    void          *recvbuf,  
    int           recvcnt,  
    MPI_Datatype  recvtype,  
    int           root,  
    MPI_Comm      comm  
);
```



Collective communication



```
int MPI_Scatter (  
    void          *sendbuf,  
    int           sendcnt,  
    MPI_Datatype  sendtype,  
    void          *recvbuf,  
    int           recvcnt,  
    MPI_Datatype  recvtype,  
    int           root,  
    MPI_Comm      comm  
);
```



Collective communication

To operate on messages of unequal length:

- MPI_Scatter**v**
- MPI_Gather**v**

To obtain results on all processors:

- MPI_Allreduce
- MPI_Allgather (**v**)

To Send data from all to all processes

- MPI_Alltoall(**v**)

To broadcast message

- MPI_Bcast



To synchronize between processors

- MPI_Barrier

Good programming

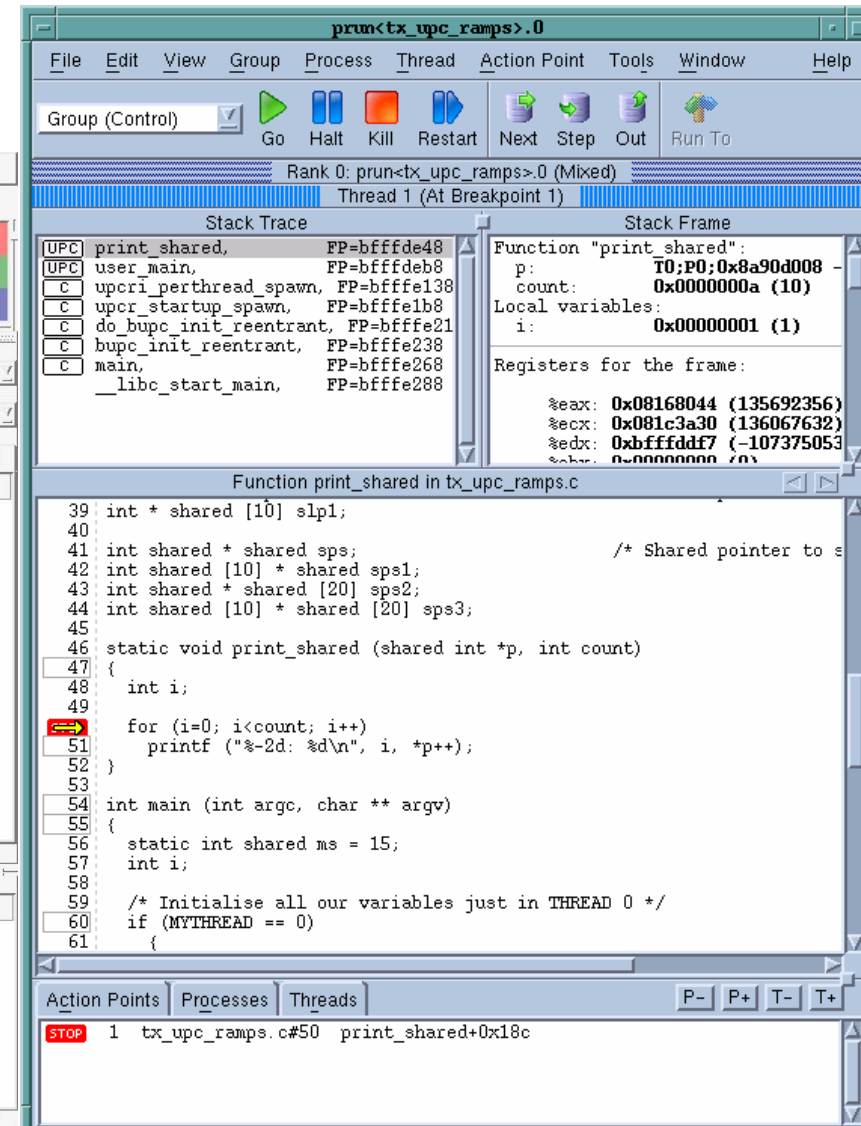
- **Reliability** – The code does not have errors and can be trusted to compute what it is supposed to compute.
- **Robustness**, which is closely related to reliability – The code has a wide range of applicability as well as the ability to detect bad data, “singular” or other problems that it can not be expected to handle, and other abnormal situations, and deal with them in a way that is satisfactory to user.
- **Portability** – The code can be transferred from one computer to another with a minimum effort and without losing reliability. Usually this means that the code has been written in a general high-level language like FORTRAN (C++) and uses no “tricks” that are dependent on the characteristic of a particular computer. Any machine characteristics that must be used are clearly delineated.
- **Maintainability** – Any code will necessary need to be changed from time to time, either to make corrections or to add enhancements, and this should be possible with minimum effort.

Code optimization through code profiling



Difference: 40 sec → 25 sec

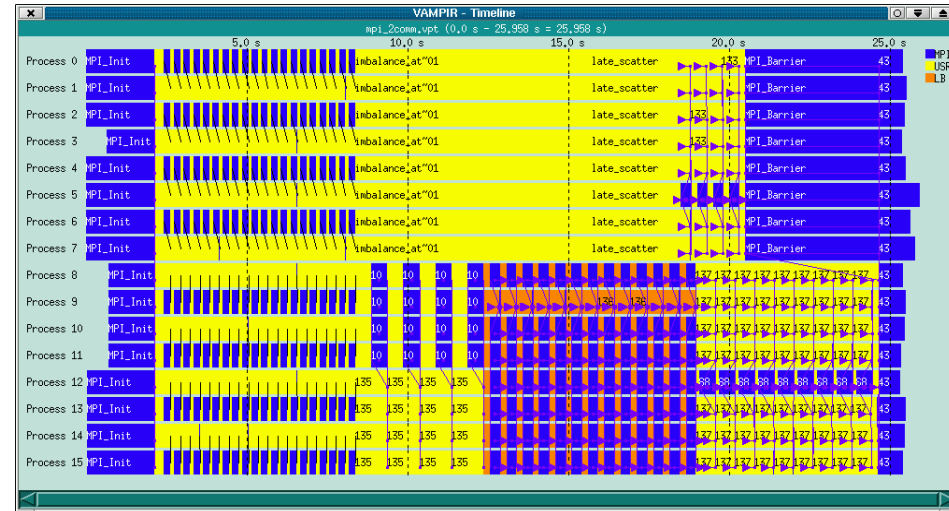
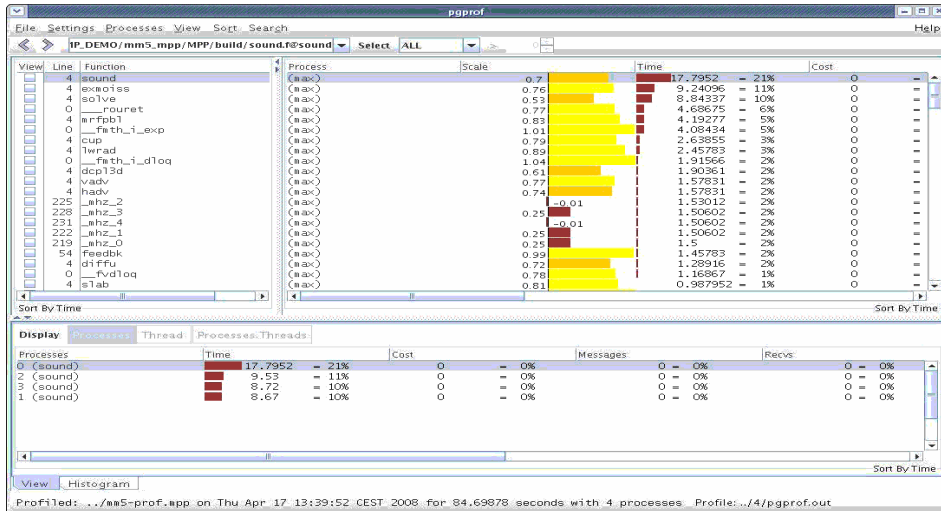
Totalview

[illegible]

Performance analysis tools

pgprof

vampire




crayPAT

IMP

100.0% | 100.0% | 512 | Total

59.8%	59.8%	306	stepfx_
17.6%	77.3%	90	getrusage
8.0%	85.4%	41	stepfy_
6.2%	91.6%	32	integr_
2.0%	93.6%	10	gradco_
1.0%	94.5%	5	__write
0.8%	95.3%	4	filerx_



How to learn MPI programming?

Just do it!